

1. A particular GPU has the following performance specifications:
 - a. It is capable of performing calculations at a rate of 2,000 GFLOPs (giga-floating point operations per second).
 - b. It can transfer 200 gigabytes of data to/from global memory per second.

If a CUDA kernel is to fully utilize the available processing power of this GPU, what's the minimum required compute to global memory access (CGMA) ratio? Assume that the size of a floating point value is 4 bytes.

Answer: 40. 50 giga-floating point values can be transferred to/from global memory per second. In order for these to be processed at 2,000 GFLOPS, the compute to access ratio should be at least $2000/50 = 40$.

2. In general, it is preferable to statically allocate arrays in GPU global memory (e.g. by declaring the array with `__device__`) rather than dynamically allocating them (e.g. with `cudaMalloc`). However, static allocation is often not possible in practice – briefly explain why this is.

Answer: For static allocation the size of the array must be known at compile time. Often, the actual data size is only known by run time, because it is dependent on user input, or the size of a file, or something else.

3. In a CUDA kernel, register spilling occurs when a local variable can't be stored in register memory and is instead stored in a special region of global memory. Although global memory is many times slower than register memory, register spilling does not always result in a big reduction in performance. Why is this?

Answer: Local variables that have been spilled out into global memory are likely to be cached locally in the streaming multiprocessor's L1 cache, which is much faster than global memory.

4. The following table shows some properties relating to the different levels of CUDA Compute Ability (also called Compute Capability):

Technical specifications	Compute ability (version)													
	1.0	1.1	1.2	1.3	2.x	3.0	3.5	3.7	5.0	5.2	5.3	6.0	6.1	
Maximum number of threads per block	512				1024									
Warp size	32													
Maximum number of resident blocks per multiprocessor	8					16			32					
Maximum number of resident warps per multiprocessor	24		32		48	64								
Maximum number of resident threads per multiprocessor	768		1024		1536	2048								

With a Compute Capability 3.0 GPU and a block dimension of (50, 10, 1), what is the maximum number of blocks that can be resident in a multiprocessor? For this question assume that the kernel doesn't make use of any shared memory.

Answer: 4. Each block contains 500 threads, so the limit according to "resident threads per MP" is $2048/500=4$. The limit according to "maximum number of blocks per multiprocessor" is 16. The minimum of these two limits is 4, so the "maximum number of resident threads per MP" is the limiting factor here.

5. The following table shows some properties relating to the different levels of CUDA Compute Ability (also called Compute Capability):

Technical specifications	Compute ability (version)														
	1.0	1.1	1.2	1.3	2.x	3.0	3.2	3.5	3.7	5.0	5.2	5.3	6.0	6.1	6.2
Maximum number of resident blocks per multiprocessor	8					16					32				
Maximum number of resident warps per multiprocessor	24		32		48	64									
Maximum number of resident threads per multiprocessor	768		1024		1536	2048									
Maximum amount of shared memory per multiprocessor	16 KB				48 KB				112 KB	64 KB	96 KB	64 KB		96 KB	64 KB
Maximum amount of shared memory per thread block	48 KB														

With a Compute Capability 6.0 GPU, a block dimension of (4, 4, 2) and a kernel that uses 1KB of shared memory per block, what is the maximum number of blocks that can be resident in a multiprocessor?

Answer: 32. Each block contains 32 threads, so the limit according to “resident threads per MP” is $2048/32=64$. The limit according to “shared memory per MP” is $64/1=64$. The limit according to “maximum number of blocks per multiprocessor” is 32. The minimum of these three limits is 32, so the “maximum number of resident blocks per MP” is the limiting factor here.

6. A CUDA streaming multiprocessor will stall if all of its resident warps of threads are in the Waiting state, meaning none of them can be executed. Describe one optimization that can be done to your code in order to reduce stalling.

Answer: Increasing occupancy to reduce the chance that all warps are Waiting. Organising memory to take advantage of coalesced memory access, reducing the amount of time spent waiting for data. Using shared memory to reduce the amount of data that must be transferred to/from global memory. If not much shared memory is required, increasing the amount of L1 cache available while reducing the size of shared memory.

7. A CUDA kernel contains the following code:

```
if (threadIdx.z == 0) {
    doSomething();
} else {
    doSomethingElse();
}
```

This kernel is launched on a grid with dimension (10, 10, 1), where the dimension of each block is (4, 4, 4). Assuming a standard warp size of 32 threads, in how many warps will execution diverge as a result of the `if` statement above? **Hint:** Recall that CUDA linearizes threads by ordering them by z co-ordinate first, then y co-ordinate, then x co-ordinate.

Answer: 100. There are 64 threads (= 2 warps) per block. In each block the first 16 threads will have `z == 0`, while the rest will have `z != 0`. This means that the first warp of each block will diverge, since it contains some threads where the `if` statement is true and some where it is false. The second warp in each block won't diverge. In total there are 100 blocks, so if one warp in each block diverges then 100 warps will diverge in total.

8. The following is a CUDA kernel that takes as input **NUM_THREAD** initial x positions and then calculates a new position for each by simulating **T** time steps (assume **T** is a **const**

unsigned int that is defined elsewhere in the code). Each initial *x* value (and thread) has a unique value for some parameter that controls how the variable will change. The kernel uses Euler's method, but you don't need to know the details of that to answer this question. The function **derivative** involves performing ten floating point operations.

```
__device__ float d_xs[NUM_THREADS];
__constant__ float d_paramValues[NUM_THREADS];

const float dt = 0.01f;

__global__ void kernel()
{
    int globalId = blockIdx.x * blockDim.x + threadIdx.x;
    float x = d_xs[globalId];
    float paramValue = d_paramValues[globalId];

    for (int i = 0; i < T; i++) {
        x += dt * derivative(x, paramValue);
    }
    d_xs[globalId] = x;
}
```

What is the compute to global memory (CGMA) ratio of this kernel? Your answer should be some expression involving the number of steps (*T*).

Answer:

Compute part: each iteration of the **for** loop involves twelve floating point operations: an addition, a multiplication, and a call to **derivative**. This loop runs **T** times, so the total number of computations will be $12 \cdot T$.

Global memory access part: each thread reads one value from global memory, and writes back one value. There are therefore two global memory accesses per thread.

Overall ratio: Dividing the number of computations by the number of global memory accesses gives $12T / 2 = 6T$.

9. In C programming it's very common to use a **struct** to store structured data. For example, the position of a particle in four dimensional space (co-ordinates *w*, *x*, *y*, *z*) could be represented by objects of the following type:

```
typedef struct { float w, x, y, z; } ParticlePosition;
```

If an application involves simulating **N** particles, an array containing the particles' positions can be allocated in GPU memory with the following code:

```
__device__ ParticlePosition particles[N];
```

If the code is written this way, memory will be organised so that the four co-ordinates associated with the first particle are stored first, then the co-ordinates of the second particle, and so on. Assuming thread *i* processes the movement of particle *i*, why might this not be the optimal way of organising memory? What arrangement could be used instead?

Answer: If adjacent threads access adjacent memory addresses simultaneously the memory accesses can be coalesced into a single transfer. Using an array of structure arrangement means that the memory addresses that nearby threads access will be spread out in memory. Instead, a structure of arrays could be used, where a different array is used to store each of the position components.

10. Although the design of GPUs is based on a single instruction multiple data (SIMD) approach, there are a number of reasons why a GPU can't be considered a purely SIMD device. Briefly describe one of these reasons.

Answer: 1) A GPU contains multiple streaming multi-processors. Each multi-processor can be executing a different instruction at once. 2) Threads within a block can all take different paths through the code (divergence). This gives the appearance of different threads in the same block executing different instructions in parallel, although in reality each branch is executed sequentially.

11. You are told that a program has a speed-up factor of 100 when executed in parallel on 20 processors. You run the program on one processor and it takes 1,500 seconds to complete. How long would you expect the program to take to run on 20 processors, and what would its efficiency be?

Answer: $S = T_S/T_P = 1500/T_P = 100 \rightarrow T_P = 15$ seconds. Efficiency $E = S/(\# \text{ processors}) = 100/20 = 5$.

12. A program takes 100 seconds to run in parallel with two processors. Ten seconds of this time is spent loading data from disk, which cannot be parallelised, while the remaining 90 seconds consist of code that parallelises perfectly with linear speed-up. How long should you expect the program to take to run on four processors? According to Amdahl's law, what's the theoretical maximum speed-up that could be achieved by adding additional processors?

Answer: If the number of processors is doubled from 2 to 4 then with linear speed-up the time for the parallel part of the program will be halved, to 45 seconds. In total the program should therefore take $45+10=55$ seconds to run on four processors. The proportion of the program that can't be parallelised is $10/100 = 0.1$, therefore according to Amdahl's law the maximum possible speed-up is $1/0.1 = 10$.

13. A common model for parallelising a relatively small program involves breaking the task up into multiple threads of execution, each of which can run on different CPU cores but within the same process. Is this an example of a shared memory or distributed memory approach?

Answer: Shared memory.


14. Simulations of spiking neural networks often only involve transmitting very small amounts of data between different processes. However, for good performance the latency between any given pair of nodes should be small (i.e. messages should arrive quickly). When considering different supercomputers to run a spiking neural network simulation on, which of the following is likely to be the most important for performance: **bisection width**, **diameter** or **valency**? Briefly explain your answer; you should mention whether a smaller value or a larger value is preferable.

Answer: A small **diameter** is likely to be most important for this application. A large diameter means a message from one node to another may potentially have to pass through many other nodes to reach its destination, resulting in high latency.

15. The diameter of a fat tree network is $2 \log_2 \frac{N+1}{2}$. If you have a fat tree network where it takes 0.5ms to send a message from one node to another, and there are $N=1000$ nodes, what is the maximum expected delay (to the nearest 0.1ms) between a pair of nodes? **Note:** in the actual class test you won't have to calculate \log_2 .

Answer: Diameter = $2 \log_2(500.5) = 17.9$. Maximum Latency = $17.9 \times 0.5\text{ms} = 9.0\text{ms}$

16. Another type of network that we haven't studied is a 2D square network. This is similar to a 2D torus but without the "wrap around" connections at the edges. The diagram below shows a 2D square network with the number of nodes $N=16$.



What is the **bisection width** of this specific network? What will the **valency** be for any 2D square network in general?

Answer: The bisection width is four, since cutting the network in down the middle will divide it into two equally sized halves, and this will involve cutting four connections. The valency for any 2D square network will also be four, since this is how many connections each node makes.